

```
import keras
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D
from keras.layers import Dense, Dropout, Flatten
from keras.optimizers import Adam
import matplotlib.pyplot as plt
import seaborn as sns
import cv2
import pickle
import pandas as pd
import numpy as np
import random
```

```
!git clone https://bitbucket.org/jadslim/german-traffic-signs
```



```
Cloning into 'german-traffic-signs'...
remote: Enumerating objects: 6, done.
remote: Counting objects: 100% (6/6), done.
remote: Compressing objects: 100% (6/6), done.
remote: Total 6 (delta 0), reused 0 (delta 0), pack-reused 0
Unpacking objects: 100% (6/6), 117.80 MiB | 5.29 MiB/s, done.
Updating files: 100% (4/4), done.
```

```
!ls german-traffic-signs
```

```
signnames.csv  test.p  train.p  valid.p
```

```
df = pd.read_csv("german-traffic-signs/signnames.csv")
```

```
df.head()
```

	ClassId	SignName	
0	0	Speed limit (20km/h)	
1	1	Speed limit (30km/h)	
2	2	Speed limit (50km/h)	
3	3	Speed limit (60km/h)	
4	4	Speed limit (70km/h)	

Next steps:

[Generate code with df](#)

[View recommended plots](#)

```
with open("german-traffic-signs/train.p", mode = "rb") as training:
    train = pickle.load(training)

with open("german-traffic-signs/valid.p", mode = "rb") as validation:
    valid = pickle.load(validation)

with open("german-traffic-signs/test.p", mode = "rb") as test:
    test = pickle.load(test)
```


```
X_train, y_train = train["features"], train["labels"]
X_validation, y_validation = valid["features"], valid["labels"]
X_test, y_test = test["features"], test["labels"]
```

```
print(X_train.shape)
print(X_validation.shape)
print(X_test.shape)
```

```
(34799, 32, 32, 3)
(4410, 32, 32, 3)
(12630, 32, 32, 3)
```

```
index = np.random.randint(1, len(X_train))
plt.imshow(X_train[index])
print("Image Label: = {}".format(y_train[index]))
```

Image Label: = 18



```
#Convert images to Grayscale
#Histogram Equalization
#Normalization

#Shuffling Images
from sklearn.utils import shuffle
X_train, y_train = shuffle(X_train, y_train)

def preprocessing(img):
    #Convert to Grayscale Images
    img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    #Applying Histogram Equalization
    img = cv2.equalizeHist(img)
    #Normalization
    img = img/255
    return img

X_train_processed = np.array(list(map(preprocessing, X_train)))
X_validation_processed = np.array(list(map(preprocessing, X_validation)))
X_test_processed = np.array(list(map(preprocessing, X_test)))

#Reshape X_train, X_test, X_validation



X_train_processed = X_train_processed.reshape(X_train_processed.shape[0], X_train_processed.shape[1], X_train_processed.shape[2], 1)
X_test_processed = X_test_processed.reshape(X_test_processed.shape[0], X_test_processed.shape[1], X_test_processed.shape[2], 1)
X_validation_processed = X_validation_processed.reshape(X_validation_processed.shape[0], X_validation_processed.shape[1], X_validation_processed.shape[2], 1)

print(X_train_processed.shape)
print(X_test_processed.shape)
print(X_validation_processed.shape)

(34799, 32, 32, 1)
(12630, 32, 32, 1)
(4410, 32, 32, 1)

i = random.randint(1, len(X_train))
plt.imshow(X_train_processed[i].squeeze(), cmap = "gray")
plt.figure()
plt.imshow(X_train[i].squeeze())
```

<matplotlib.image.AxesImage at 0x7fc0c26a48b0>



```

model = Sequential()

model.add(Conv2D(filters = 32,
                  kernel_size = (5,5),
                  activation = "relu",
                  input_shape = X_train_processed.shape[1:]))

model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout(0.25)) #To Reduce Overfitting of the data

model.add(Conv2D(filters = 64,
                  kernel_size = (3,3),
                  activation = "relu"))

model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(256, activation = "relu"))
model.add(Dropout(0.5))

model.add(Dense(43, activation = "softmax"))

model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 32)	832
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
dropout (Dropout)	(None, 14, 14, 32)	0
conv2d_1 (Conv2D)	(None, 12, 12, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
dropout_1 (Dropout)	(None, 6, 6, 64)	0
flatten (Flatten)	(None, 2304)	0
dense (Dense)	(None, 256)	590080
dropout_2 (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 43)	11051

```

=====
Total params: 620459 (2.37 MB)
Trainable params: 620459 (2.37 MB)
Non-trainable params: 0 (0.00 Byte)

```

```
model.compile(optimizer = Adam(learning_rate=0.001),  
              loss = "sparse_categorical_crossentropy",  
              metrics = ["accuracy"])
```

```
history = model.fit(X_train_processed, y_train,  
                   batch_size = 128, #How many images will be feeding at once  
                   epochs = 30,  
                   validation_data = (X_validation_processed, y_validation))
```

```
Epoch 2/30  
272/272 [=====] - 2s 6ms/step - loss: 0.5591 - accuracy: 0.8264 - val_loss: 0.3328 - val_accuracy: 0.9066  
Epoch 3/30  
272/272 [=====] - 2s 6ms/step - loss: 0.3512 - accuracy: 0.8899 - val_loss: 0.2326 - val_accuracy: 0.9274  
Epoch 4/30  
272/272 [=====] - 2s 6ms/step - loss: 0.2642 - accuracy: 0.9169 - val_loss: 0.1927 - val_accuracy: 0.9376  
Epoch 5/30  
272/272 [=====] - 2s 6ms/step - loss: 0.2091 - accuracy: 0.9340 - val_loss: 0.1648 - val_accuracy: 0.9483  
Epoch 6/30  
272/272 [=====] - 2s 6ms/step - loss: 0.1751 - accuracy: 0.9447 - val_loss: 0.1638 - val_accuracy: 0.9544  
Epoch 7/30  
272/272 [=====] - 2s 6ms/step - loss: 0.1546 - accuracy: 0.9515 - val_loss: 0.1491 - val_accuracy: 0.9526  
Epoch 8/30  
272/272 [=====] - 2s 7ms/step - loss: 0.1310 - accuracy: 0.9590 - val_loss: 0.1377 - val_accuracy: 0.9544  
Epoch 9/30  
272/272 [=====] - 2s 7ms/step - loss: 0.1114 - accuracy: 0.9648 - val_loss: 0.1164 - val_accuracy: 0.9639  
Epoch 10/30  
272/272 [=====] - 2s 6ms/step - loss: 0.1045 - accuracy: 0.9669 - val_loss: 0.1277 - val_accuracy: 0.9601  
Epoch 11/30  
272/272 [=====] - 2s 6ms/step - loss: 0.0966 - accuracy: 0.9692 - val_loss: 0.1195 - val_accuracy: 0.9619  
Epoch 12/30  
272/272 [=====] - 2s 6ms/step - loss: 0.0862 - accuracy: 0.9723 - val_loss: 0.0967 - val_accuracy: 0.9728  
Epoch 13/30  
272/272 [=====] - 2s 6ms/step - loss: 0.0783 - accuracy: 0.9752 - val_loss: 0.1388 - val_accuracy: 0.9605  
Epoch 14/30  
272/272 [=====] - 2s 6ms/step - loss: 0.0722 - accuracy: 0.9762 - val_loss: 0.1232 - val_accuracy: 0.9637  
Epoch 15/30  
272/272 [=====] - 2s 6ms/step - loss: 0.0706 - accuracy: 0.9769 - val_loss: 0.1176 - val_accuracy: 0.9635  
Epoch 16/30  
272/272 [=====] - 2s 8ms/step - loss: 0.0631 - accuracy: 0.9794 - val_loss: 0.1037 - val_accuracy: 0.9696  
Epoch 17/30  
272/272 [=====] - 2s 7ms/step - loss: 0.0589 - accuracy: 0.9809 - val_loss: 0.1154 - val_accuracy: 0.9680  
Epoch 18/30  
272/272 [=====] - 2s 6ms/step - loss: 0.0587 - accuracy: 0.9805 - val_loss: 0.1015 - val_accuracy: 0.9712  
Epoch 19/30  
272/272 [=====] - 2s 6ms/step - loss: 0.0515 - accuracy: 0.9827 - val_loss: 0.1092 - val_accuracy: 0.9676  
Epoch 20/30  
272/272 [=====] - 2s 6ms/step - loss: 0.0518 - accuracy: 0.9824 - val_loss: 0.0965 - val_accuracy: 0.9737  
Epoch 21/30  
272/272 [=====] - 2s 6ms/step - loss: 0.0519 - accuracy: 0.9830 - val_loss: 0.0915 - val_accuracy: 0.9741  
Epoch 22/30  
272/272 [=====] - 2s 6ms/step - loss: 0.0507 - accuracy: 0.9834 - val_loss: 0.1027 - val_accuracy: 0.9717  
Epoch 23/30  
272/272 [=====] - 2s 6ms/step - loss: 0.0454 - accuracy: 0.9848 - val_loss: 0.0984 - val_accuracy: 0.9721  
Epoch 24/30  
272/272 [=====] - 2s 7ms/step - loss: 0.0470 - accuracy: 0.9845 - val_loss: 0.0837 - val_accuracy: 0.9769  
Epoch 25/30  
272/272 [=====] - 2s 6ms/step - loss: 0.0413 - accuracy: 0.9869 - val_loss: 0.1003 - val_accuracy: 0.9728  
Epoch 26/30
```

```

272/272 [=====] - 2s 6ms/step - loss: 0.0423 - accuracy: 0.9859 - val_loss: 0.1250 - val_accuracy: 0.9662
Epoch 28/30
272/272 [=====] - 2s 6ms/step - loss: 0.0398 - accuracy: 0.9863 - val_loss: 0.1342 - val_accuracy: 0.9635
Epoch 29/30
272/272 [=====] - 2s 6ms/step - loss: 0.0392 - accuracy: 0.9872 - val_loss: 0.1207 - val_accuracy: 0.9689
Epoch 30/30
272/272 [=====] - 2s 6ms/step - loss: 0.0368 - accuracy: 0.9875 - val_loss: 0.1262 - val_accuracy: 0.9680

```



#Create loss and accuracy plot

```

import matplotlib.pyplot as plt
sns.set()
plt.figure(figsize=(20,5))
plt.subplot(1,2,1)
plt.plot(history.history['loss'], color='b', label="Training Loss")
plt.plot(history.history['val_loss'], color='r', label="Validation Loss")
plt.legend()
plt.xlabel("Number of Epochs")
plt.ylabel("Loss")

plt.subplot(1,2,2)
plt.plot(history.history['accuracy'], color='b', label="Training Accuracy")
plt.plot(history.history['val_accuracy'], color='r', label="Validation Accuracy")
plt.legend()
plt.xlabel("Number of Epochs")
plt.ylabel("Accuracy")
plt.show()

```



```

score = model.evaluate(X_test_processed, y_test, verbose = 0)
print("Test Loss: ", score[0])
print("Test Accuracy: ", score[1])


```

```
Test Loss: 0.1725344955921173  
Test Accuracy: 0.9600158333778381
```


```
history.history.keys()
```

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
plt.plot(history.history['loss'])  
plt.plot(history.history['val_loss'])  
plt.title('model loss')  
plt.ylabel('loss')  
plt.xlabel('epoch')  
plt.legend(['train', 'val'], loc='upper left')  
plt.show()
```



```
plt.plot(history.history['accuracy'])  
plt.plot(history.history['val_accuracy'])  
plt.title('model accuracy')  
plt.ylabel('accuracy')  
plt.xlabel('epoch')  
plt.legend(['train', 'val'], loc='upper left')  
plt.show()
```

```
prediction = model.predict(X_test_processed)
```

```
395/395 [=====] - 1s 3ms/step
```

```
#Generate confusion matrix
```

```
import numpy as np
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Get predictions from the model
predictions = model.predict(X_test_processed)
```

```
# Generate the confusion matrix
matrix = confusion_matrix(y_test, np.argmax(predictions, axis=1))
```

```
# Plot the confusion matrix
plt.figure(figsize=(12, 7))
sns.heatmap(matrix, annot=True)
plt.show()
```

395/395 [=====] - 1s 2ms/step

